

Chapter 8 Applied Exercises

To start with a clean session, we first remove all packages except the base R packages (we will load anything needed later) and all variables.


```
rm(list = ls())
pkgs = names(sessionInfo())$otherPkgs
if (!is.null(pkgs)) {
  detach_list <- paste0("package:", pkgs)
  lapply(detach_list, FUN = detach, character.only = TRUE)
}
```

```
## [[1]]
## NULL
##
## [[2]]
## NULL
##
## [[3]]
## NULL
##
## [[4]]
## NULL
##
## [[5]]
## NULL
##
## [[6]]
## NULL
##
## [[7]]
## NULL
##
## [[8]]
## NULL
##
## [[9]]
## NULL
##
## [[10]]
## NULL
##
## [[11]]
## NULL
##
## [[12]]
## NULL
##
## [[13]]
## NULL
##
## [[14]]
## NULL
```

```
##
## [[15]]
## NULL
##
## [[16]]
## NULL
```

Exercise 3

```
p = seq(0,1, length = 100)
# Gini Index
gini= 2*p*(1-p)
err = apply(cbind(1-p, p), FUN = min, MARGIN = 1)
entropy = -p*log(p) - (1-p)*log(1-p)
error.mat = cbind(gini, err, entropy)
matplot(p, error.mat, xlab = expression(italic(hat(p)[mk])), type = "l", lwd = 2, col = 1:3, lty = 1:3)
legend("topright", legend = c("Gini", "Classification", "Entropy"), lwd = 2, col = 1:3, lty = 1:3)
```



To see how to write mathematical expressions on a plot, type `?plotmath`. Pairwise minimum command `pmin` could be used above instead of the `apply()` function.

Exercise 7

To see not only how changes in the number of trees and m affect the test error, but also how they interact, we depict them on the same plot. We find the test MSE as a function of the number of trees for values $m = p, p/2, \sqrt{p}$.

```
library(randomForest)

## randomForest 4.6-12
## Type rfNews() to see new features/changes/bug fixes.

library(MASS)
# define test and train
set.seed(1)
train = sample(1:nrow(Boston), size = nrow(Boston)/2)
test = -train
test.y = Boston$medv[test]
# Grids for ntree and mtry
(ntree_seq = seq(1, 1000, by = 20))


## [1] 1 21 41 61 81 101 121 141 161 181 201 221 241 261 281 301 321
## [18] 341 361 381 401 421 441 461 481 501 521 541 561 581 601 621 641 661
## [35] 681 701 721 741 761 781 801 821 841 861 881 901 921 941 961 981

(mtry_seq = c(13, round(13/2), round(sqrt(13))))

## [1] 13 6 4

err_mat = matrix(NA, nrow = length(ntree_seq), ncol = length(mtry_seq),
                 dimnames = list(NULL, NULL))
# error estimates for the grid
for (i_tree in 1:length(ntree_seq)) {
  for (j_mtry in 1:length(mtry_seq)) {
    rf_boston = randomForest(medv ~ ., data = Boston, ntree = ntree_seq[i_tree],
                             mtry = mtry_seq[j_mtry], subset = train)
    rf_pred = predict(rf_boston, newdata = Boston[test, ])
    err_mat[i_tree, j_mtry] = mean((rf_pred - test.y)^2)
  }
}

matplot(ntree_seq, err_mat, type = "l", lty = 1, col = c("blue", "darkgreen", "brown"), lwd = 1, ylim =
legend("topright", legend = paste0("m = ", mtry_seq), col = c("blue", "darkgreen", "brown"), lwd = 1)
```



Using 400 trees is sufficient to give good performance and $m \simeq p/2 = 6$ leads to the best test MSE. Random forest outperforms bagging in this example.

The 3D plot shows that a very small m would result in large error, and the error falls as we increase m . This plot is, however, difficult to interpret. So we make an attempt toward drawing 2D plots which could reflect the interaction between the two variable of interest.

In order to use the package `ggplot2`, we need to tidy the data first, i.e. make each row of the data represent one point.

Suppose we have a grid $x = \{x_i\}_{i=1}^n$ on the X-axis and $y = \{y_i\}_{i=1}^q$ on the y-axis and a matrix $Z_{n \times q} = [z_{ij}]_{ij}$ where $z_{ij} = f(x_i, y_j)$. The goal is to represent each point as a single row, i.e. make a one-to-one correspondence between the grids, i.e. between vectors x and y , and also between the grids and the matrix Z . To achieve this, we create a matrix $X_{n \times q}$ out of vector x by copying x to each column of X . We also create Y by copying the vector y to each row of the matrix $Y_{n \times q}$.

Now let x be the number of splits and y equal to m , the number of predictors considered in each split. We have

```
X = matrix(ntree_seq, nrow = length(ntree_seq), ncol = length(mtry_seq))
Y = matrix(mtry_seq, nrow = length(ntree_seq), ncol = length(mtry_seq), byrow = TRUE)
head(X)
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12] [,13]
## [1,]    1    1    1    1    1    1    1    1    1    1    1    1    1
## [2,]    3    3    3    3    3    3    3    3    3    3    3    3    3
## [3,]    5    5    5    5    5    5    5    5    5    5    5    5    5
## [4,]    7    7    7    7    7    7    7    7    7    7    7    7    7
## [5,]    9    9    9    9    9    9    9    9    9    9    9    9    9
## [6,]   11   11   11   11   11   11   11   11   11   11   11   11   11
```

```
head(Y)
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12] [,13]
## [1,]    1    2    3    4    5    6    7    8    9    10    11    12    13
## [2,]    1    2    3    4    5    6    7    8    9    10    11    12    13
## [3,]    1    2    3    4    5    6    7    8    9    10    11    12    13
## [4,]    1    2    3    4    5    6    7    8    9    10    11    12    13
## [5,]    1    2    3    4    5    6    7    8    9    10    11    12    13
## [6,]    1    2    3    4    5    6    7    8    9    10    11    12    13
```

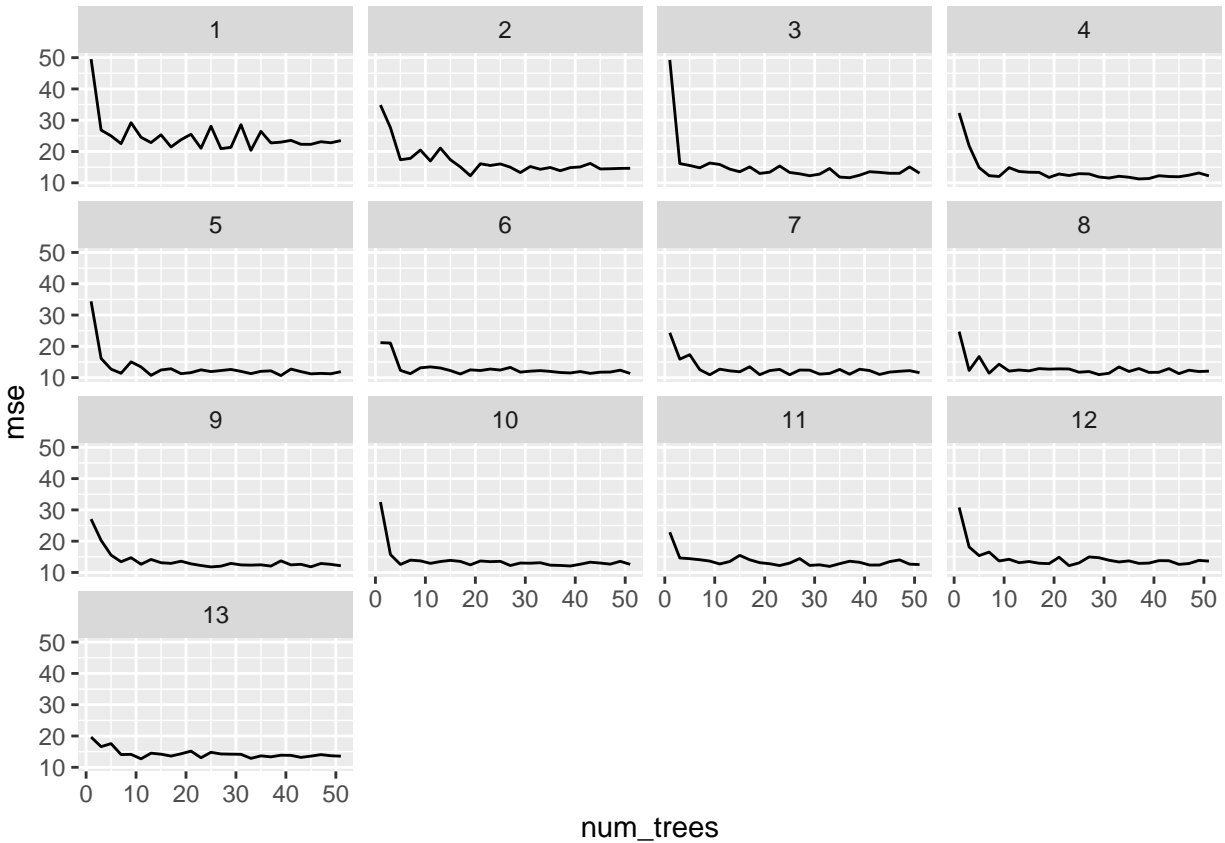
Note the use of `byrow = TRUE`. Now that all elements of matrices `X`, `Y` and `err_mat` correspond, we generate a data frame by converting all these matrices to columns:

```
df_tree = data.frame(num_trees = c(X), num_vars = c(Y), mse = c(err_mat))
head(df)
```

```
##
## 1 function (x, df1, df2, ncp, log = FALSE)
## 2 {
## 3     if (missing(ncp))
## 4         .Call(C_df, x, df1, df2, log)
## 5     else .Call(C_dnf, x, df1, df2, ncp, log)
## 6 }
```

We can use the data frame `df_tree` in `ggplot2` to draw different plots.

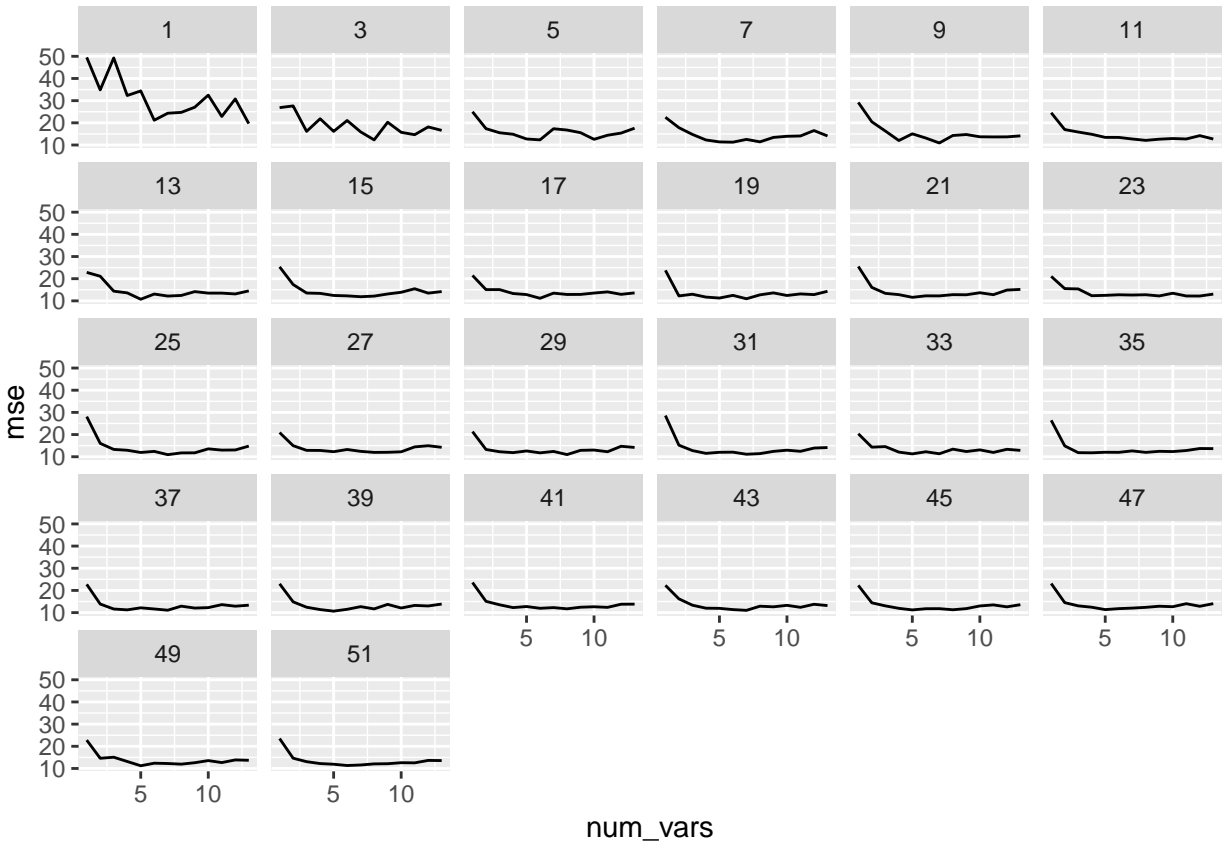
```
library(ggplot2)
ggplot(data = df_tree) +
  geom_line(mapping = aes(x = num_trees, y = mse)) +
  facet_wrap(~num_vars)
```



Using the number of variables as the facet, we could see how the number of trees affect the error estimate. By increasing the number of trees to 5, we see a considerable improvement in accuracy for different levels of m , and the improvement slows down for larger number of trees.

We might be tempted to state that increasing the number of variables does not have much effect on reducing the error for m greater than 4, implying that $m = 4$ might be a good choice. However, note that this is not the best plot for making this conclusion. The y-axis scale is too large to be able to distinguish the differences between the error values for large enough values of `num_trees` and `num_vars`. To see how the plot above might hide such information, consider `num_trees` as the facet:

```
ggplot(data = df_tree) +
  geom_line(mapping = aes(x = num_vars, y = mse)) +
  facet_wrap(~num_trees)
```



Using the number of trees as the facet, we see that the error has a downward trend and stabilizes at about 5 variables, given the values of `num_trees`. We can further adjust the scale on the y-axis to see the trend. We could also explore how changing the random seed might affect our conclusion, or investigate the effect of increasing bootstrap samples, but we will not pursue either of these issues here.

->

Exercise 8

Part 8.a)

We will predict `Sales` in the `Carseats` data, treating it as a quantitative variable.

```
library(ISLR)
library(tree)
# split the data
set.seed(1)
train = sample(1:nrow(Carseats), size = nrow(Carseats)/2)
test = -train
test.y = Carseats$Sales[test]
```

Part 8.b)

```
# fit regression tree
tree.carseats = tree(Sales ~ ., data = Carseats, subset = train)
tree.carseats
```

```
## node), split, n, deviance, yval
##      * denotes terminal node
##
## 1) root 200 1526.000  7.335
##    2) ShelfLoc: Bad,Medium 164 1084.000  6.815
##      4) Price < 120.5 103  541.000  7.776
##        8) Age < 50.5 39  133.100  9.278
##          16) Price < 104.5 25  83.050  9.808
##            32) ShelfLoc: Bad 7  22.340  8.284 *
##            33) ShelfLoc: Medium 18  38.140 10.400 *
##          17) Price > 104.5 14  30.510  8.333
##            34) Advertising < 3.5 5  6.584  6.922 *
##            35) Advertising > 3.5 9  8.444  9.117 *
##        9) Age > 50.5 64  266.200  6.860
##          18) Price < 92 17  67.560  8.628
##            36) Income < 85 9  23.830  7.512 *
##            37) Income > 85 8  19.940  9.882 *
##          19) Price > 92 47  126.300  6.221
##            38) ShelfLoc: Bad 16  29.890  5.059 *
##            39) ShelfLoc: Medium 31  63.660  6.820
##              78) CompPrice < 107 5  4.229  5.206 *
##              79) CompPrice > 107 26  43.900  7.131 *
##        5) Price > 120.5 61  287.000  5.192
##          10) Age < 66.5 49  181.400  5.654
##            20) CompPrice < 148 37  129.600  5.208
##              40) Advertising < 10.5 25  75.180  4.627 *
##              41) Advertising > 10.5 12  28.370  6.418 *
##            21) CompPrice > 148 12  21.680  7.031 *
##          11) Age > 66.5 12  52.410  3.303
##            22) Price < 132 5  16.300  4.780 *
##            23) Price > 132 7  17.410  2.249 *
##    3) ShelfLoc: Good 36  196.000  9.706
##      6) Price < 113 10  19.960 12.080 *
##      7) Price > 113 26  97.810  8.792
##        14) Price < 142.5 20  54.100  9.337
##          28) CompPrice < 133 15  24.220  8.778 *
##          29) CompPrice > 133 5  11.130 11.010 *
##        15) Price > 142.5 6  17.930  6.973 *
```

What if we use Gini index for growing the tree?

```
tree.carseats_gini = tree(Sales ~ ., data = Carseats, subset = train, split = "gini")
tree.carseats_gini
```

```
## node), split, n, deviance, yval
##      * denotes terminal node
##
## 1) root 200 1526 7.335 *
```

It leads to no partitioning! Hence, it seems best to use other packages such as **rpart** for splitting according to the Gini index. It is also not clear to me the default splitting according to deviance is equivalent to cross-entropy. A little algebra shows that deviance contribution of each region R_m should be normalized by

As a result of these points, although the package `tree` could be a strating point, it is not my choice for final results in a regression tree. However, we continue to use the package `tree` and the default for its argument `split` which is `deviance`.

[illegible]

Similar to the classification tree considered in the chapter's lab, shelving location seems to be the most important variable. The first branch differentiates the **Good** locations from **Medium** and **Bad** locations. Price also appears to be important, as many branches distinguish levels of price.

```
# predict for test observations
tree_preds = predict(tree.carseats, newdata = Carseats[test, ])
# test MSE:
mse_tree = mean((tree_preds - test.y)^2)
mse_tree
```

9


Part 8.c)

Now we prune the tree to see whether we can improve prediction accuracy.

```
set.seed(1)
cv_carseats = cv.tree(tree.carseats)
cv_carseats

## $size
## [1] 18 17 16 15 14 12 11 10 9 8 7 6 5 4 3 1
##
## $dev
## [1] 1127.510 1163.270 1163.270 1169.097 1169.097 1153.636 1129.237
## [8] 1139.639 1139.232 1109.839 1135.897 1126.203 1218.821 1205.007
## [15] 1331.927 1612.664
##
## $k
## [1] -Inf 15.48181 15.53599 18.69038 18.74886 21.05038 23.79480
## [8] 25.78579 26.01210 30.10435 32.74801 53.28569 72.33061 78.19599
## [15] 141.73781 251.22901
##
## $method
## [1] "deviance"
##
## attr(,"class")
## [1] "prune" "tree.sequence"

plot(cv_carseats$size, cv_carseats$dev, type = "b")
```



The cross-validation MSE estimates is minimized for 8 splits.

```
# find the pruned tree
prune_carseats = prune.tree(tree.carseats, best = 8)
# predict for the pruned tree
prune_preds = predict(prune_carseats, newdata = Carseats[test, ])
# test mse for the pruned tree
mse_prune_carseats = mean((prune_preds - test.y)^2)
mse_prune_carseats
```

```
## [1] 5.09085
```

It results in larger test MSE, so seems not to improve accuracy (for this specific test set).

Exercise 8.d)

```
library(randomForest)

## randomForest 4.6-12
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
## The following object is masked from 'package:ggplot2':
##
##     margin
```

```

set.seed(1)
bag_carseats = randomForest(Sales ~ ., data = Carseats, subset = train,
                             mtry = ncol(Carseats)-1, ntree = 500, importance = TRUE)
preds_bag_carseats = predict(bag_carseats, newdata = Carseats[test, ])
mse_bag_carseats = mean((preds_bag_carseats - test.y)^2)
mse_bag_carseats

```

```
## [1] 2.554292
```

Using bagging would improve the prediction accuracy.

Note that we set `mtry` equal to number of all the predictors, which is supposed to give us the bagging estimator, which is a special case of random forest estimator. However, it appears that doing so would not exactly result in a bagging estimator, as introduced in the textbook. We would expect the bagging estimator to be certain, since there is no randomness when all predictors are considered in each split. However, it appears that the algorithm for `randomForest` generates other sources of randomness. That is why we have set the seed in the code chunk above.

If we do not set the random seed, we would get different results each time we use `randomForest()` even for maximum value of `mtry` and `ntree` being equal to one.


Now we determine which variables are most important:

```
importance(bag_carseats)
```

```
##           %IncMSE IncNodePurity
## CompPrice  14.032030    129.568747
## Income     5.523038     75.448682
## Advertising 13.571285    131.246840
## Population  1.968853     63.042648
## Price      56.863812    504.158108
## ShelfLoc   44.720455    323.055042
## Age        22.225468    194.915976
## Education  4.823966     40.810991
## Urban      -1.902185      8.746566
## US         6.632887    14.599565
```

```
varImpPlot(bag_carseats)
```

bag_carseats



The variables Price and ShelfLoc appear to be the most important variables according to both measures, one based on contribution of the variable to the out-of-bag estimate of the test error (on the LHS) and the other based on contribution of the variable to the total node impurity in the training data (on the RHS). The two importance measures generally yield similar results in terms of relative importance of the variables in determining sales. An exception is the variable US which is more important according to the OOB estimates.

Part 8.e)

```
set.seed(1)
rf_carseats <- randomForest(Sales ~ ., data = Carseats, subset = train,
                             ntree = 500, importance = TRUE)
preds_rf_carseats <- predict(rf_carseats, newdata = Carseats[test, ])
mse_rf_carseats <- mean((preds_rf_carseats - test.y)^2)
mse_rf_carseats
```

```
## [1] 3.30763
```


```
importance(rf_carseats)
```

```
##           %IncMSE IncNodePurity
## CompPrice   9.849043    134.17665
## Income      3.534622    130.84360
## Advertising 13.075334    139.40128
## Population  -0.612195    100.34668
## Price       35.530402    380.27956
```

```
## ShelfLoc 31.015873 234.62966
## Age 16.680174 202.18673
## Education 2.563741 68.75977
## Urban -1.569224 16.99182
## US 6.400241 31.12594
```

```
varImpPlot(rf_carseats)
```


rf_carseats



The random forest estimator has larger estimate for test MSE, but leads to similar relative importance for variables. The optimal value of m would depend on the bias-variance trade-off: the lower the number of variables considered in each split, the higher the variance, but the lower the correlation between variables and as a result, the variance of the estimator.

Below, we see how the prediction accuracy changes with changes in m :

```
set.seed(1)
Mse_rf_carseats <- rep(NA, length = 10)
for (i in 1:10) {
  rf_temp <- randomForest(Sales ~ ., data = Carseats, subset = train,
                           mtry = i, ntree = 500, importance = TRUE)
  preds_temp = predict(rf_temp, newdata = Carseats[test, ])
  Mse_rf_carseats[i] = mean((preds_temp - test.y)^2)
}
plot(1:10, Mse_rf_carseats, xlab = "m", type = "b")
```



The bagging estimator appears to have the lowest (estimate of) test MSE among different random forest estimators. Hence, the decrease in variance due to reducing m is not enough to offset the larger bias resulted by considering fewer variables in each split.

This might be related to the idea of why we would use small values of m when we have many highly correlated variables (mentioned in the textbook). On one hand, we find m to be the largest here and on the other hand, we have a small number of observations and variables, and the variables are not highly correlated.

Exercise 9

```
rm(list = ls())
library(ISLR)
# knowing the data
dim(OJ) # 1070*18

## [1] 1070 18

sum(is.na(OJ))

## [1] 0

summary(OJ) # Purchase and Store7 factor variables
```

```
## Purchase WeekofPurchase StoreID PriceCH PriceMM
## CH:653 Min. :227.0 Min. :1.00 Min. :1.690 Min. :1.690
## MM:417 1st Qu.:240.0 1st Qu.:2.00 1st Qu.:1.790 1st Qu.:1.990
## Median :257.0 Median :3.00 Median :1.860 Median :2.090
## Mean :254.4 Mean :3.96 Mean :1.867 Mean :2.085
```

```
##          3rd Qu.:268.0  3rd Qu.:7.00  3rd Qu.:1.990  3rd Qu.:2.180
##          Max.   :278.0  Max.   :7.00  Max.   :2.090  Max.   :2.290
##          DiscCH          DiscMM          SpecialCH          SpecialMM
## Min.   :0.00000  Min.   :0.0000  Min.   :0.0000  Min.   :0.0000
## 1st Qu.:0.00000  1st Qu.:0.0000  1st Qu.:0.0000  1st Qu.:0.0000
## Median :0.00000  Median :0.0000  Median :0.0000  Median :0.0000
## Mean   :0.05186  Mean   :0.1234  Mean   :0.1477  Mean   :0.1617
## 3rd Qu.:0.00000  3rd Qu.:0.2300  3rd Qu.:0.0000  3rd Qu.:0.0000
## Max.   :0.50000  Max.   :0.8000  Max.   :1.0000  Max.   :1.0000
##          LoyalCH          SalePriceMM          SalePriceCH          PriceDiff
## Min.   :0.000011  Min.   :1.190  Min.   :1.390  Min.   : -0.6700
## 1st Qu.:0.325257  1st Qu.:1.690  1st Qu.:1.750  1st Qu.: 0.0000
## Median :0.600000  Median :2.090  Median :1.860  Median : 0.2300
## Mean   :0.565782  Mean   :1.962  Mean   :1.816  Mean   : 0.1465
## 3rd Qu.:0.850873  3rd Qu.:2.130  3rd Qu.:1.890  3rd Qu.: 0.3200
## Max.   :0.999947  Max.   :2.290  Max.   :2.090  Max.   : 0.6400
## Store7          PctDiscMM          PctDiscCH          ListPriceDiff
## No :714  Min.   :0.0000  Min.   :0.00000  Min.   :0.000
## Yes:356  1st Qu.:0.0000  1st Qu.:0.00000  1st Qu.:0.140
##          Median :0.0000  Median :0.00000  Median :0.240
##          Mean   :0.0593  Mean   :0.02731  Mean   :0.218
##          3rd Qu.:0.1127  3rd Qu.:0.00000  3rd Qu.:0.300
##          Max.   :0.4020  Max.   :0.25269  Max.   :0.440
##          STORE
## Min.   :0.000
## 1st Qu.:0.000
## Median :2.000
## Mean   :1.631
## 3rd Qu.:3.000
## Max.   :4.000
```

```
# knowing the response
contrasts(OJ$Purchase) # MM is 1 and CH is 0; MM
```

```
##      MM
## CH   0
## MM   1
```

Part 9.a)

```
train = sample(1:nrow(OJ), size = 800)
test = -train
test_purchase = OJ$Purchase[test]
```

Part 9.b)

```
library(tree)
tree_oj <- tree(Purchase ~ ., data = OJ, subset = train)
summary(tree_oj)
```

```
##
## Classification tree:
```



```
## tree(formula = Purchase ~ ., data = OJ, subset = train)
## Variables actually used in tree construction:
## [1] "LoyalCH"      "SalePriceMM"  "SpecialCH"    "WeekofPurchase"
## [5] "ListPriceDiff" "PctDiscMM"
## Number of terminal nodes: 9
## Residual mean deviance: 0.7141 = 564.8 / 791
## Misclassification error rate: 0.1512 = 121 / 800
```

Deviance is a measure of fit which is equivalent to RSS for cases such as least squares. The smaller the deviance, the better the fit (to the training data). The training error rate is 0.15125 and the tree has 9 terminal nodes.

Part 9.c)

Here is a detailed text summary of the table:


```
tree_oj
```

```
## node), split, n, deviance, yval, (yprob)
##      * denotes terminal node
##
## 1) root 800 1072.00 CH ( 0.60750 0.39250 )
##    2) LoyalCH < 0.48285 305 318.60 MM ( 0.21639 0.78361 )
##      4) LoyalCH < 0.136698 103 45.76 MM ( 0.05825 0.94175 ) *
##      5) LoyalCH > 0.136698 202 245.80 MM ( 0.29703 0.70297 )
##        10) SalePriceMM < 2.04 115 112.30 MM ( 0.19130 0.80870 )
##          20) SpecialCH < 0.5 102 81.59 MM ( 0.13725 0.86275 ) *
##          21) SpecialCH > 0.5 13 17.32 CH ( 0.61538 0.38462 ) *
##        11) SalePriceMM > 2.04 87 119.20 MM ( 0.43678 0.56322 )
##          22) WeekofPurchase < 246.5 21 17.22 MM ( 0.14286 0.85714 ) *
##          23) WeekofPurchase > 246.5 66 91.25 CH ( 0.53030 0.46970 ) *
##    3) LoyalCH > 0.48285 495 421.10 CH ( 0.84848 0.15152 )
##      6) LoyalCH < 0.764572 236 277.80 CH ( 0.72458 0.27542 )
##        12) ListPriceDiff < 0.235 91 126.10 MM ( 0.49451 0.50549 )
##          24) PctDiscMM < 0.196197 72 97.80 CH ( 0.58333 0.41667 ) *
##          25) PctDiscMM > 0.196197 19 16.57 MM ( 0.15789 0.84211 ) *
##      13) ListPriceDiff > 0.235 145 112.60 CH ( 0.86897 0.13103 ) *
##      7) LoyalCH > 0.764572 259 84.69 CH ( 0.96139 0.03861 ) *
```

Branches that lead to terminal nodes are indicated by astrisk symbols. In each row, we can see the split criterion, and the number of observations, the deviance and the overall prediction for the branch. The first number in parentheses is the fraction of observations that take on the value MM and the second is the fraction that take on the value CH.

Part 9.d)

```
plot(tree_oj)
text(tree_oj, pretty = FALSE, cex = 0.7)
```



According to the tree we grew, brand loyalty to CH is the most important predictor for sales. Given `LoyalCH` is greater than 0.75 or smaller than 0.027, other variables seem to play little role in determining **Purchase**. But among customers who do not have much loyalty to either brand, price difference does influence purchases, and if price of MM is low enough, compared to CH, they tend to buy MM. Loyalty matters even among these medium-loyalty customers: it takes a much cheaper MM (`PriceDiff < -0.165`) to persuade someone a bit loyal to CH ($0.5 < \text{LoyalCH} < 0.75$) to buy MM (compared to `PriceDiff < 0.05` for $0.28 < \text{LoyalCH} < 0.5$).

Part 9.e)

```

preds_oj = predict(tree_oj, newdata = OJ[test, ], type = "class")
table_oj <- table(test_purchase, preds_oj)
table_oj

```

```

##           preds_oj
## test_purchase  CH  MM
##           CH 155  12
##           MM  46  57

```

Test error rate is about 19%:

```

1 - sum(diag(table_oj))/sum(table_oj)

```

```

## [1] 0.2148148

```

```

mean(test_purchase != preds_oj)

```

```

## [1] 0.2148148

```

Part 9.f)

```
set.seed(1)
(cv_oj = cv.tree(tree_oj, FUN = prune.misclass))


## $size
## [1] 9 5 2 1
##
## $dev
## [1] 152 150 149 314
##
## $k
## [1]      -Inf    1.750000    4.333333 173.000000
##
## $method
## [1] "misclass"
##
## attr("class")
## [1] "prune"          "tree.sequence"
```

k represents the tuning parameter in the minimization problem solved in cost-complexity pruning. When k is equal to ∞ , it is as if we are maximizing $|T|$ which gives us the tree that we start with, the one with $|T_0|$ terminal nodes. k is represented as α in ISLR.

In the example above, the sequence of trees T_α above, found by solving the cost-complexity minimization, are a subset of trees resulted by weakest link pruning. Although weakest link pruning yields all subsets of the original tree (i.e. sub-trees with sizes equal to 8, 7, 6, 5, 4, 3, 2, 1), only a subset of them (trees with size equal to 9, 5, 2, 1) solve the cost-complexity problem.

Part 9.g)

```
plot(cv_oj$size, cv_oj$dev, type = "b")
```



Part 9.h)


The error rate (confusingly labelled above as **dev** in the output) is minimized for a tree with 2 terminal nodes. But the difference in test error rates for trees of size 2, 5 and 8 is very small.

Part 9.i)

All the trees we found above by using cost-complexity pruning (done by `cv.tree`) are a subset of sub-trees that can be found using weakest link pruning (done by `prune.tree`). Hence, any tree with a given size $|T|$ is a tree found by weakest link pruning (successively ommiting branches) if and only if it is found by cost-complexity pruning (solving the optimization problem).

As a result, given any tree, all we need to know about the the optimal sub-tree is the size of it. Without knowing anything about the shape of the optimal sub-tree (beyond its size), we can easily find it through weakest link pruning.

```
prune_oj = prune.tree(tree_oj, best = 2)
plot(prune_oj)
text(prune_oj, pretty = FALSE, cex = 0.9)
```



Part 9.j)

The training and test error rates should be the same, since the predicted values for trees with size 6 and 8 are the same. The split that is omitted by pruning led to the same predicted value for both leaves (it only existed since it increased node purity).

We already showed that the classification test error rates are the same for the full tree and the pruned one. Below, we confirm the training error rates are also the same for them:

```
# training error rate for full tree
preds_full <- predict(tree_oj, newdata = OJ[train, ], type = "class")
mean(OJ$Purchase[train] != preds_full)
```

```
## [1] 0.15125
```

```
# training error rate for optimal tree
preds_prune <- predict(prune_oj, newdata = OJ[train, ], type = "class")
mean(OJ$Purchase[train] != preds_prune)
```

```
## [1] 0.17625
```

The full tree has a smaller test error rate, which is not strange, given very similar C.V. error rates between the full tree and the pruned one.

Part 9.k)

See part 9.j.

Exercise 10

Knowing the data:

```
library(ISLR)
# know the data
dim(Hitters)
```

```
## [1] 322 20
```

```
summary(Hitters) # League, Division and NewLeague are quantitative
```

```
##      AtBat      Hits      HmRun      Runs
##  Min.   : 16.0   Min.    :  1   Min.    : 0.00   Min.    :  0.00
## 1st Qu.:255.2   1st Qu.: 64   1st Qu.: 4.00   1st Qu.: 30.25
## Median :379.5   Median : 96   Median : 8.00   Median : 48.00
## Mean   :380.9   Mean    :101   Mean    :10.77   Mean    : 50.91
## 3rd Qu.:512.0   3rd Qu.:137   3rd Qu.:16.00   3rd Qu.: 69.00
## Max.   :687.0   Max.    :238   Max.    :40.00   Max.    :130.00
##
##      RBI      Walks      Years      CAtBat
##  Min.   :  0.00   Min.    :  0.00   Min.    : 1.000   Min.    : 19.0
## 1st Qu.: 28.00   1st Qu.: 22.00   1st Qu.: 4.000   1st Qu.: 816.8
## Median : 44.00   Median : 35.00   Median : 6.000   Median :1928.0
## Mean   : 48.03   Mean     :38.74   Mean     : 7.444   Mean     :2648.7
## 3rd Qu.: 64.75   3rd Qu.: 53.00   3rd Qu.:11.000   3rd Qu.:3924.2
## Max.   :121.00   Max.     :105.00   Max.     :24.000   Max.     :14053.0
##
##      CHits      CHmRun      CRuns      CRBI
##  Min.   :  4.0   Min.    :  0.00   Min.    :  1.0   Min.    :  0.00
## 1st Qu.: 209.0   1st Qu.: 14.00   1st Qu.: 100.2   1st Qu.: 88.75
## Median : 508.0   Median : 37.50   Median : 247.0   Median :220.50
## Mean   : 717.6   Mean     :69.49   Mean     :358.8   Mean     :330.12
## 3rd Qu.:1059.2   3rd Qu.: 90.00   3rd Qu.: 526.2   3rd Qu.:426.25
## Max.   :4256.0   Max.     :548.00   Max.     :2165.0   Max.     :1659.00
##
##      CWalks      League Division      PutOuts      Assists
##  Min.   :  0.00   A:175   E:157   Min.    :  0.0   Min.    :  0.0
## 1st Qu.: 67.25   N:147   W:165   1st Qu.: 109.2   1st Qu.:  7.0
## Median : 170.50                      Median : 212.0   Median : 39.5
## Mean   : 260.24                      Mean    : 288.9   Mean    :106.9
## 3rd Qu.: 339.25                      3rd Qu.: 325.0   3rd Qu.:166.0
## Max.   :1566.00                      Max.     :1378.0   Max.     :492.0
##
##      Errors      Salary      NewLeague
##  Min.   :  0.00   Min.    : 67.5   A:176
## 1st Qu.:  3.00   1st Qu.:190.0   N:146
## Median :  6.00   Median :425.0
## Mean   :  8.04   Mean    :535.9
## 3rd Qu.:11.00   3rd Qu.:750.0
## Max.   :32.00   Max.    :2460.0
##
##      NA's      :59
```

```
sum(is.na(Hitters)) # 59 missing
```


```
## [1] 59
```

```
# know the dependent variable: annual salary in thousands of dollars (see ?Hitters)
sum(is.na(Hitters$Salary)) # 59 missing
```

```
## [1] 59
```

```
hist(Hitters$Salary) # very skewed, long tail
```

Histogram of Hitters\$Salary



Here, we know the missing values arise from `Hitters$Salary`. Nonetheless, in general for seeing what variables have missing values, we can also run the code below:

```
length(Hitters)
```

```
## [1] 20
```

```
sapply(Hitters, FUN = function(x) sum(is.na(x)))
```

```
##      AtBat      Hits      HmRun      Runs      RBI      Walks      Years
##         0         0         0         0         0         0         0
##    CAtBat    CHits    CHmRun    CRuns    CRBI    CWalks    League
##         0         0         0         0         0         0         0
## Division PutOuts Assists  Errors  Salary NewLeague
##         0         0         0         0        59         0
```

Part 10.a)

Cleaning the data:

```
# we omit the missing, as suggested in the exercise:
Hits <- na.omit(Hitters)
# log-transform salary:
Hits$lsalary <- log(Hits$Salary)
hist(Hits$lsalary)
```



The log-transformation alleviates the skewness to a great extent. From now on, we work with `Hits` as the data and `lsalary` as the response.

Part 10.b)

```
train = 1:200
test = -train
```

Part 10.c)

We want to draw training MSE as a function of the shrinkage parameter.

```
# grid for lambda
(lambda_grid <- 10^seq(0, -4, length = 9))
```

```
## [1] 1.0000000000 0.3162277660 0.1000000000 0.0316227766 0.0100000000
## [6] 0.0031622777 0.0010000000 0.0003162278 0.0001000000
```



```

# training MSE
library(gbm)

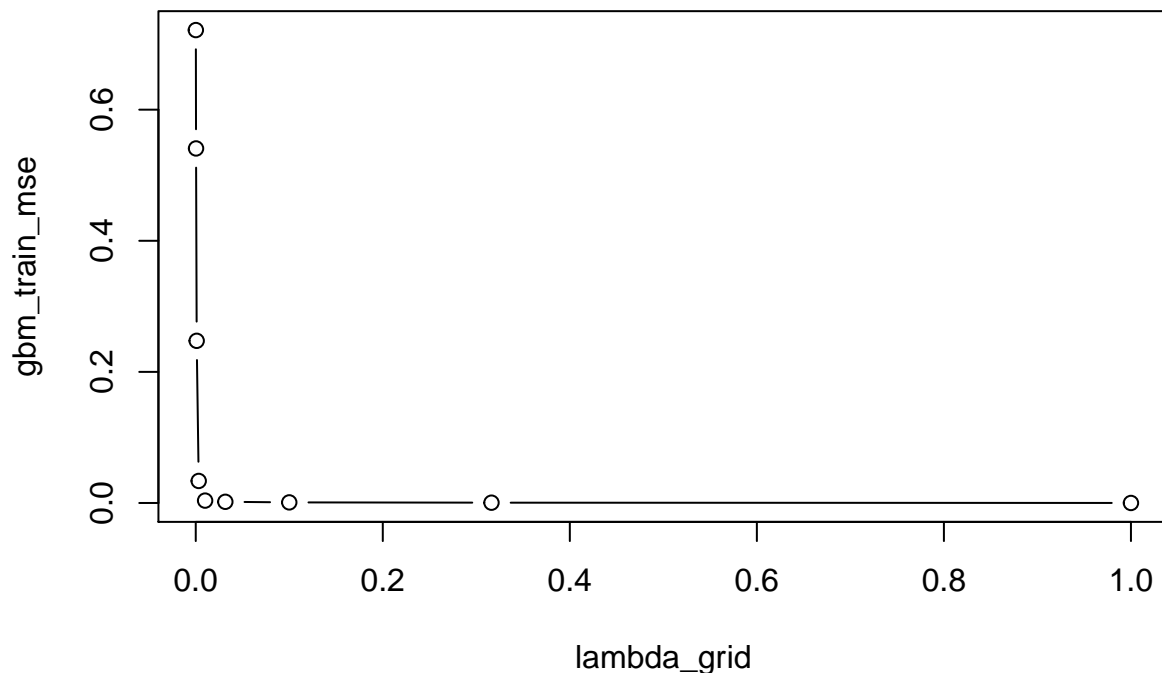
## Loading required package: survival
## Loading required package: lattice
## Loaded gbm 2.1.3

library(ggplot2)
set.seed(1)
gbm_fits <- vector(mode = "list", length = length(lambda_grid))
gbm_train_preds <- vector(mode = "list", length = length(lambda_grid))
gbm_train_mse <- vector(length = length(lambda_grid))
for (i in 1:length(lambda_grid)) {
  gbm_fits[[i]] <- gbm(lsalary ~ ., data = Hits[train, ], distribution = "gaussian",
    n.trees = 1000, interaction.depth = 1, shrinkage = lambda_grid[i])
  gbm_train_preds[[i]] <- predict(gbm_fits[[i]], newdata = Hits[train, ], n.trees = 1000)
  gbm_train_mse[i] <- mean((Hits$lsalary[train] - gbm_train_preds[[i]])^2)
}

```


Here is the plot:

```
plot(lambda_grid, gbm_train_mse, type = "b")
```



As we see above the training MSE goes down to zero very quickly as we decrease λ . The direction of change in the training MSE could be seen more clearly by using a log-log scale:

```
plot(lambda_grid, gbm_train_mse, log = "xy", type = "b")
```



The log-log transformation shows the percentage change in training MSE as we change λ . Training MSE goes down as we increase λ . What the above plots show is that we get a better fit to the training data as we decrease the rate of learning. But the result is probably due to fixing the number of trees; for larger values of λ we would need fewer trees to avoid overfitting. But since we are fixing the number of trees, the decline in the training MSE is likely to be due to overfitting for large enough shrinkage parameters. We need to compute the test MSE for finding the optimal value of λ .

As we saw, in this example it was easier to use the base plotting function in R for our purpose. Nevertheless, `ggplot` enables the implementation of much more elaborate details. For example, see here. ->

Part 10.d)

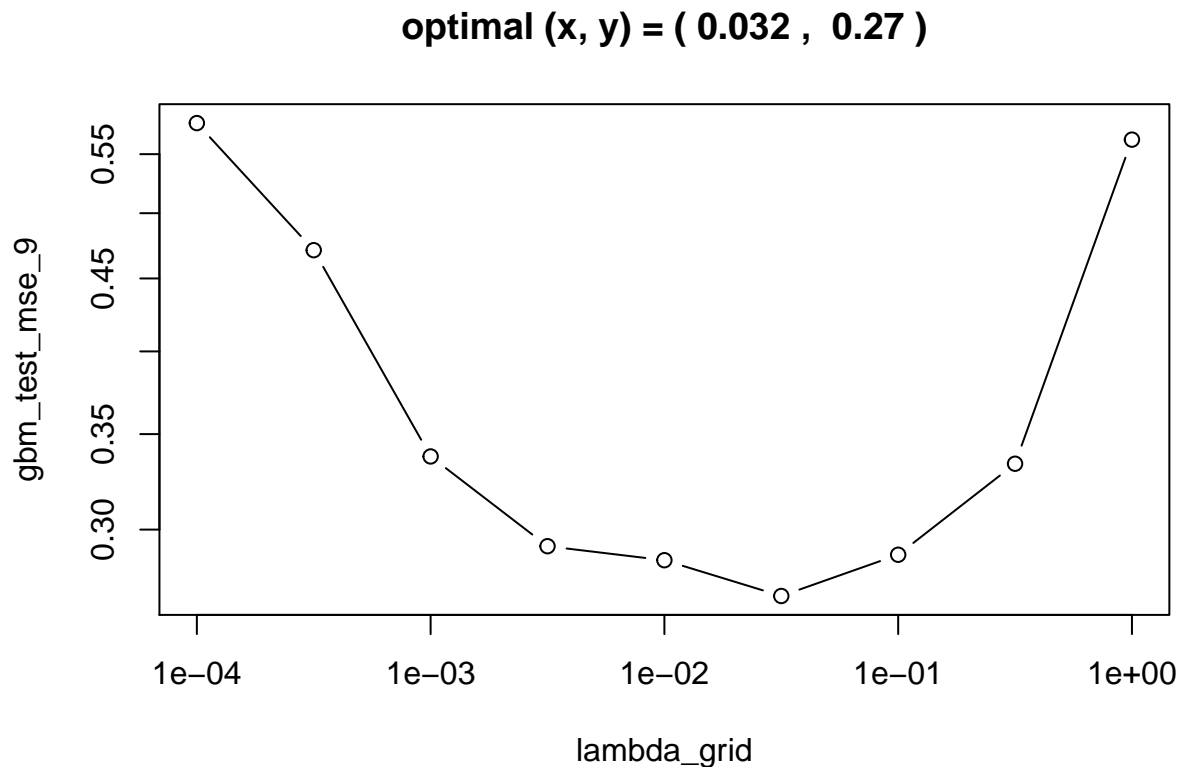
To enable working with different combinations, we write a function:

```
test_mse_ex10 <-
function(data, lambda_grid, n.trees) {
  set.seed(1)
  gbm_test_mse <- rep(NA, length = length(lambda_grid))
  for (i in 1:length(lambda_grid)) {
    gbm_fit <- gbm(lsalary ~ . - Salary - lsalary, data = data[train, ], distribution = "gaussian",
                  shrinkage = lambda_grid[i], interaction.depth = 1, n.trees = n.trees)
    gbm_preds <- predict(gbm_fit, newdata = data[test, ], n.tree = n.trees)
    gbm_test_mse[i] <- mean((data$lsalary[test] - gbm_preds)^2)
  }
}
```

```
gbm_test_mse
}
```

We use the same grid for the test MSE.


```
gbm_test_mse_9 <- test_mse_ex10(Hits, lambda_grid = lambda_grid, n.trees = 1000)
opt_mse_9 <- c(lambda_grid[gbm_test_mse_9 == min(gbm_test_mse_9)],
               min(gbm_test_mse_9))
plot(lambda_grid, gbm_test_mse_9, log = "xy", type = "b")
title(paste("optimal (x, y) = (", round(opt_mse_9[1], 3), ", ", round(opt_mse_9[2], 3), ")"))
```



We can prove that the minimum in the log-log plot is the same as the minimum in the non-transformed plot, since x and y are positive (it can be shown using the fact that $\frac{d \log(y)}{d \log(x)} = \frac{dy}{dx} \times \frac{x}{y}$, so whenever the derivative of the non-transformed function changes sign from negative to positive, so does the derivative of transformed). Since we do not have many grid points, the optimal point might not be the minimum point in the plot above. Below, we use a finer grid, but only increasing the number of grid points to 20, but also narrowing our focus on a smaller region:

```
lambda_grid_20 <- 10^seq(from = log(0.02), to = log(0.7), length = 20)
gbm_test_mse_20 <- test_mse_ex10(Hits, lambda_grid = lambda_grid_20, n.trees = 1000)
opt_mse_20 <- c(lambda_grid_20[gbm_test_mse_20 == min(gbm_test_mse_20)],
                min(gbm_test_mse_20))
plot(lambda_grid_20, gbm_test_mse_20, log = "xy", type = "b")
title(paste("optimal (x, y) = (", round(opt_mse_20[1], 3), ", ", round(opt_mse_20[2], 3), ")"))
```

optimal (x, y) = (0.051 , 0.258)



```
c(lambda_grid_20[gbm_test_mse_20 == min(gbm_test_mse_20)], min(gbm_test_mse_20))
```

```
## [1] 0.05101609 0.25788970
```


An interesting point is that the test MSE seems to be sensitive to the amount of λ , in the sense that a small change in λ may casue a relatively large change in the test MSE. This is better seen using a finer grid, which we will investigate in the next subsection. However, we will use the amount of λ we find here in the next rest of the exercise.

Digression: the effect of finder grids

The picture below shows the change in the plot if we only increase the number of grid points from 20 to 100, given the same interval:

```
lambda_grid_100 <- 10^seq(from = log(0.02), to = log(0.7), length = 100)
gbm_test_mse_100 <- test_mse_ex10(Hits, lambda.grid = lambda_grid_100, n.trees = 1000)
opt_mse <- c(lambda_grid_100[gbm_test_mse_100 == min(gbm_test_mse_100)],
             min(gbm_test_mse_100))
plot(lambda_grid_100, gbm_test_mse_100, log = "xy", type = "b")
title(paste("optimal (x, y) = (", round(opt_mse[1], 3), ", ", round(opt_mse[2], 3), ")"))
```

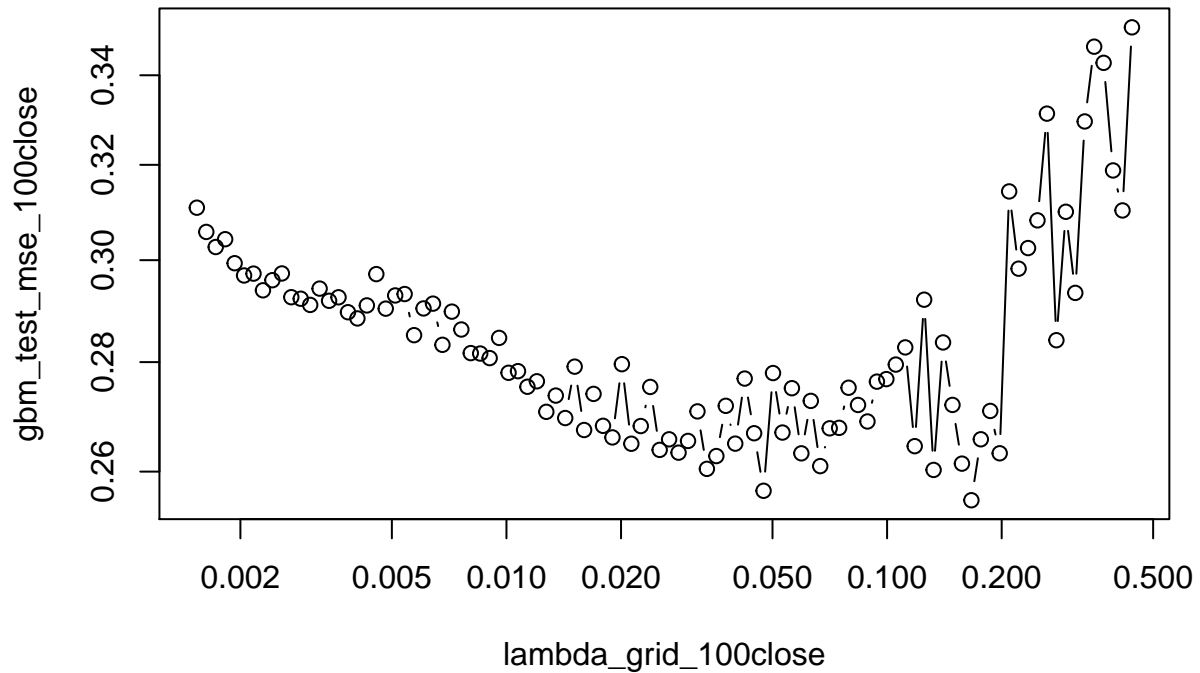
optimal (x, y) = (0.117 , 0.248)



Below we look even closer. Even a closer look shows us that the optimal λ should be between 0.02 and 0.2. But, due to the fluctuations, it does not give us much information beyond that. So we keep the number of grid points equal to 100, but narrow the interval:

```
lambda_grid_100close <- 10^seq(from = log(0.06), to = log(0.7), length = 100)
gbm_test_mse_100close <-
  test_mse_ex10(Hits, lambda.grid = lambda_grid_100close, n.trees = 1000)
opt_mse <- c(lambda_grid_100close[gbm_test_mse_100close == min(gbm_test_mse_100close)],
             min(gbm_test_mse_100close))
plot(lambda_grid_100close, gbm_test_mse_100close, log = "xy", type = "b")
title(paste("optimal (x, y) = (", round(opt_mse[1], 3), ", ", round(opt_mse[2], 3), ")"))
```


optimal (x, y) = (0.167 , 0.255)



How would the optimal λ change when we increase the number of trees? Does it improve the test MSE? We keep the narrowest interval we experimented with, as the previous plot, and increase the number of trees from 1000 to 5000:

```
lambda_grid_100close <- 10^seq(from = log(0.06), to = log(0.7), length = 100)
gbm_test_mse_100close <- test_mse_ex10(Hits, lambda.grid = lambda_grid_100close, n.trees = 5000)
opt_mse <- c(lambda_grid_100close[gbm_test_mse_100close == min(gbm_test_mse_100close)], min(gbm_test_mse_100close))
plot(lambda_grid_100close, gbm_test_mse_100close, log = "xy", type = "b")
title(paste("optimal (x, y) = (", round(opt_mse[1], 3), ", ", round(opt_mse[2], 3), ")"))
```

optimal (x, y) = (0.015 , 0.256)



In this example, optimal λ decreases when we increase the number of trees, which is in line with the idea that when we have more trees, we can have slower pace of learning.

The test MSE is very close in all cases we studied above, except for the case where we had only 9 grid points. Hence, in this example, we would have been fine as long as we used a moderate number of grid points. However, note that this is a very small dataset and we are using only 63 observations for the test data.

Part 10.e)

Chapter 3 covered linear regression and chapter 6 covered subset selection, shrinkage and dimension reduction methods. The questions asks for combining linear regression with either of the methods in chapter 6. We will work only with linear variables, so will not make any higher-order transformation.

Best subset selection

We use cross-validation with 5 folds.

```
best_size1 <- which.min(Cv_glm)
```

The warnings arise from `cv.glm()` and the fact that it find the number of observations too small. This is a problem we cannot address unless we have more data.


To understand why this is wrong, suppose we were using validation set approach instead of cross-validation. This is like using both training and test data for finding the best size and then refitting the model on the training data and evaluating on the test data. The problem is that we are using the test data for the model

selection, which introduces overfitting. Hence, when using cross-validation, the test data in each fold could not be used in any way for training.

Second method

->

```
library(leaps)
set.seed(1)
## define the predict function for regsubsets:
predict.regsubsets =
  function (object, newdata ,id, ...) {
    form = as.formula(object$call[[2]])
    mat = model.matrix(form, newdata)
    coefi = coef(object, id = id)
    xvars = names(coefi)
    mat[,xvars]%%coefi
  }
## Compute CV error
# define folds
folds <- sample(1:5, size = length(train), replace = TRUE)
# to avoid confusion of training set with different training data in cv
Hits_cv <- Hits[train, ]
Cv_mat <- matrix(NA, nrow = 19, ncol = 5)
for (j in 1:5) {
  regfit_j <- regsubsets(lsalary ~ . - Salary, data = Hits_cv[folds != j, ], nvmax = 19)
  for (i in 1:19) {
    preds_ij <- predict(regfit_j, Hits_cv[folds == j, ], id = i)
    Cv_mat[i, j] <- mean((preds_ij - Hits_cv$lsalary[folds == j])^2)
  }
}
Cv_regs <- apply(Cv_mat, FUN = mean, MARGIN = 1)
plot(1:19, Cv_regs, type = "b")
```

The plot above implies that the optimal size is 4, i.e. with four variables. Now, we should retrieve the best model of size 4 using full training data (so far we have only used 4 folds out of 5 as training data in cross-validation):

```
regfit_full <- regsubsets(lsalary ~ . - Salary, data = Hits[train, ],
                          nvmax = 4) # four variables is enough
preds2 <- predict(regfit_full, newdata = Hits[test, ], id = 4)
(mse_best2 <- mean((preds2 - Hits$lsalary[test])^2))
```

```
## [1] 0.4942052
```

Ridge regression

```
library(glmnet)
```

```
## Loading required package: Matrix
```

```
## Loading required package: foreach
```

```
## foreach: simple, scalable parallel programming from Revolution Analytics
```

```
## Use Revolution R for scalability, fault tolerance and more.
```

```
## http://www.revolutionanalytics.com
```

```
## Loaded glmnet 2.0-10
```

```
Hits_x = model.matrix(lsalary ~ . - Salary, data = Hits)[, -1]
```

```
lambda_grid <- 10^seq(10, -2, length = 100)
```

```
fit_ridge <- glmnet(x = Hits_x[train, ], y = Hits$lsalary[train], alpha = 0,
```

```

lambda = lambda_grid, thresh = 1e-12)
cv_ridge <- cv.glmnet(x = Hits_x[test, ], y = Hits$lsalary[test], alpha = 0)
best_lam <- cv_ridge$lambda.min
preds_ridge <- predict(fit_ridge, s = best_lam, newx = Hits_x[test, ])
(mse_ridge <- mean((preds_ridge - Hits$lsalary[test])^2))

```

```
## [1] 0.4424325
```

Lasso regression

```

library(glmnet)
Hits_x = model.matrix(lsalary ~ . - Salary, data = Hits)[, -1]
lambda_grid <- 10^seq(10, -2, length = 100)
fit_lasso <- glmnet(x = Hits_x[train, ], y = Hits$lsalary[train], alpha = 1,
lambda = lambda_grid, thresh = 1e-12)
cv_lasso <- cv.glmnet(x = Hits_x[test, ], y = Hits$lsalary[test], alpha = 1)
best_lam <- cv_lasso$lambda.min
preds_lasso <- predict(fit_lasso, s = best_lam, newx = Hits_x[test, ])
(mse_lasso <- mean((preds_lasso - Hits$lsalary[test])^2))

```

```
## [1] 0.4398225
```

PCR

```

library(pls)

##
## Attaching package: 'pls'

## The following object is masked from 'package:stats':
##
## loadings

pcr_fit <- pcr(lsalary ~ . - Salary, data = Hits[train, ],
scale = TRUE, validation = "CV")
summary(pcr_fit)

```

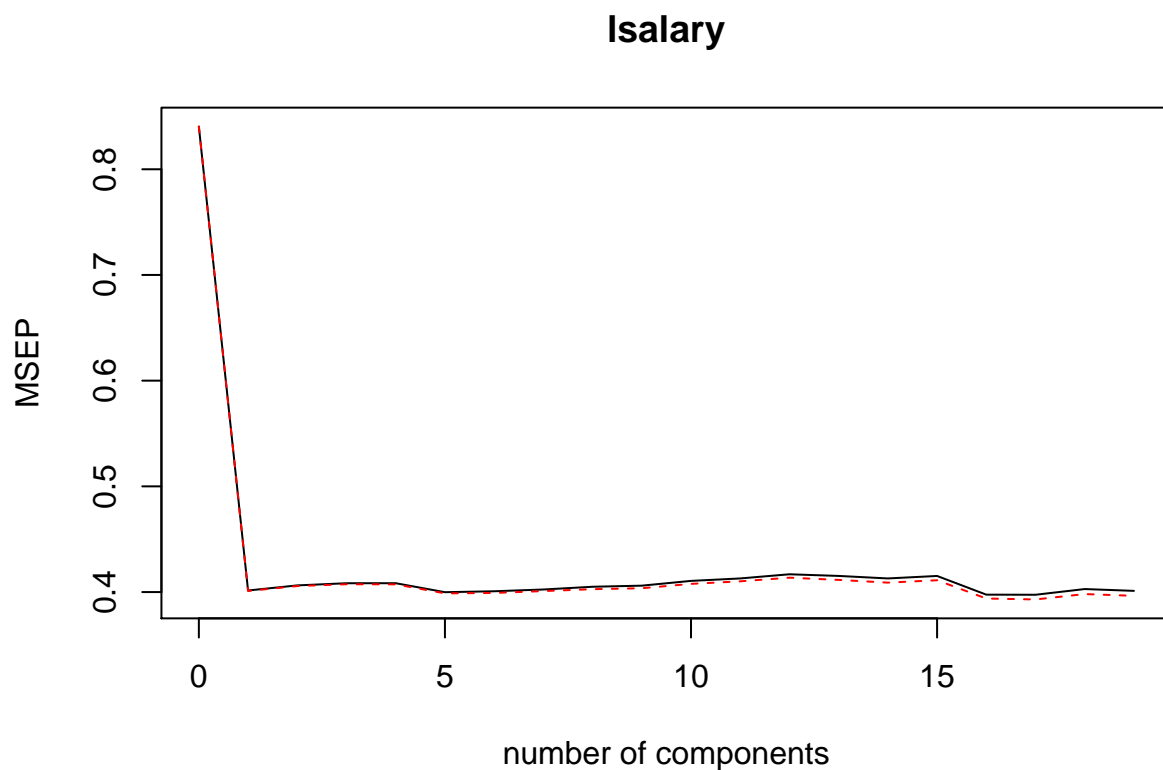
```

## Data:      X dimension: 200 19
## Y dimension: 200 1
## Fit method: svdpc
## Number of components considered: 19
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV           0.9167  0.6336  0.6374  0.6390  0.6390  0.6324  0.6330
## adjCV        0.9167  0.6333  0.6369  0.6383  0.6382  0.6315  0.6318
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV           0.6344  0.6364  0.6372  0.6407  0.6426  0.6456  0.6444
## adjCV        0.6332  0.6346  0.6353  0.6385  0.6404  0.6431  0.6415
##      14 comps 15 comps 16 comps 17 comps 18 comps 19 comps
## CV           0.6426  0.6444  0.6305  0.6304  0.6347  0.6333
## adjCV        0.6395  0.6412  0.6276  0.6269  0.6310  0.6296

```

```
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps
## X      39.40   60.67   70.76   79.64   84.73   89.29   92.38
## lsalary 52.42   52.67   53.22   53.56   54.56   55.34   55.70
##      8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## X      95.02   96.32   97.26   98.06   98.72   99.23
## lsalary 56.48   56.68   56.77   56.80   57.41   57.91
##     14 comps 15 comps 16 comps 17 comps 18 comps 19 comps
## X      99.54   99.77   99.90   99.97   99.99  100.00
## lsalary 58.68   58.75   60.27   61.27   61.27   61.49
```

```
validationplot(pcr_fit, val.type = "MSEP")
```



Choosing 1 component gives a reasonable cross-validation error.

```
preds_pcr <- predict(pcr_fit, newdata = Hits[test, ], ncomp = 1)
(cv_pcr <- mean((preds_pcr - Hits$lsalary[test])^2))
```


```
## [1] 0.4661183
```

The performance of boosting is impressive: it leads to test MSE estimate of about 0.26, while the best performance we get from the combination of a linear model with methods of chapter 6 is more than 0.43. This is despite the fact that we use an additive boosting model (we use stumps). This is an indication of nonlinearity in the data.

Part 10.f)

We use the value of λ we found with 20 grids, and use the full data:

```
best_lam <- lambda_grid_20[gbm_test_mse_20 == min(gbm_test_mse_20)]
gbm_fit <- gbm(lsalary ~ . - Salary - lsalary, data = Hits, distribution = "gaussian",
  shrinkage = best_lam, interaction.depth = 1, n.trees = 1000)
summary(gbm_fit)
```



```
##      var    rel.inf
## CATBat  CATBat 17.0952907
## CRBI    CRBI  11.0769797
## CHits   CHits  10.5285551
## CWalks  CWalks  6.5644198
## CRuns   CRuns  6.3014537
## Years   Years  5.5972824
## Walks    Walks  5.3090311
## Hits     Hits   5.2856063
## CHmRun   CHmRun  5.2416846
## PutOuts  PutOuts  4.8958056
## RBI      RBI    4.2241857
## HmRun    HmRun   4.1355981
## Runs     Runs   3.9128950
## AtBat    AtBat   3.7021136
## Errors   Errors  2.7403018
## Assists  Assists  2.2520485
## League   League  0.4980133
```

```
## Division    Division  0.3436503
## NewLeague   NewLeague  0.2950848
```

CatBat appears to be the most important variable. Perhaps more interestingly, the variables that measure values during the career seem to be much more important than the same-season measurements.

Part 10.g)

The bagging test MSE:

```
library(randomForest)
rf_fit <- randomForest(lsalary ~ . - Salary, data = Hits[train, ],
                      mtry = 19, ntree = 500, importance = TRUE)
rf_preds <- predict(rf_fit, newdata = Hits[test, ])
(rf_mse <- mean((rf_preds - Hits$lsalary[test])^2))
```

```
## [1] 0.22887
```

The test MSE is slightly smaller than the one we found for boosting. This might be because of using an additive model for boosting, while the bagging estimator is more flexible.

Exercise 11

```
library(ISLR)
# the data
dim(Caravan) # 86 variables
```

```
## [1] 5822  86
```

```
sum(is.na(Caravan))
```

```
## [1] 0
```

```
# response
summary(Caravan$Purchase)
```

```
##    No  Yes
## 5474 348
```

```
contrasts(Caravan$Purchase)
```

```
##      Yes
## No      0
## Yes     1
```

gbm() with a qualitative variable as response requires transformation to a 0-1 dummy variable:

```
Caravan$Purchase <- ifelse(Caravan$Purchase == "Yes", 1, 0)
summary(Caravan$Purchase)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.00000 0.00000 0.00000 0.05977 0.00000 1.00000
```

Part 11.a)

```
train = 1:1000
test = -train
```

The test set is nearly 4 times larger than the training set.


Part 11.b)

```
library(gbm)
set.seed(1)
gbm_fit <- gbm(Purchase ~ ., data = Caravan[train, ], distribution = "bernoulli",
               interaction.depth = 2, n.trees = 1000, shrinkage = 0.01)
```

```
## Warning in gbm.fit(x, y, offset = offset, distribution = distribution, w =
## w, : variable 50: PVRAAUT has no variation.
```

```
## Warning in gbm.fit(x, y, offset = offset, distribution = distribution, w =
## w, : variable 71: AVRAAUT has no variation.
```

```
summary(gbm_fit)
```



```
##          var    rel.inf
## PPERSAUT PPERSAUT 9.88196755
## MKOOPKLA MKOOPKLA 7.09057539
## MOPLHOOG MOPLHOOG 5.97383721
```

```

## PBRAND      PBRAND 5.42795341
## MGODGE      MGODGE 5.37533300
## MBERMIDD    MBERMIDD 3.94847387
## MINK3045    MINK3045 3.51763123
## MGODPR      MGODPR 3.25916272
## MOSTYPE     MOSTYPE 2.81168103
## MBERARBG    MBERARBG 2.52290762
## ABRAND      ABRAND 2.48318472
## MAUT2       MAUT2 2.43784324
## MSKA        MSKA 2.39996975
## MAUT1       MAUT1 2.32463277
## MSKC        MSKC 2.25575413
## PWAPART     PWAPART 2.15817221
## MSKB1       MSKB1 2.01054800
## MINK7512    MINK7512 1.86610130
## MFGEKIND    MFGEKIND 1.71887449
## MFWEKIND    MFWEKIND 1.63896537
## MRELGE      MRELGE 1.61226288
## MBERHOOG    MBERHOOG 1.56162484
## MGODOV      MGODOV 1.47899843
## MINKGEM     MINKGEM 1.44042382
## MZFONDS     MZFONDS 1.36297740
## MOPLMIDD    MOPLMIDD 1.29452540
## MGODRK      MGODRK 1.28364889
## APERSAUT    APERSAUT 1.24368912
## MINK4575    MINK4575 1.21344684
## MRELOV      MRELOV 1.19368544
## MAUTO       MAUTO 1.14689401
## MHHUUR      MHHUUR 1.14420398
## MGEMLEEF    MGEMLEEF 1.02514716
## MZPART      MZPART 1.00085325
## MHKOOP      MHKOOP 0.99160310
## MINKM30     MINKM30 0.96171242
## MOSHOOFD    MOSHOOFD 0.96138547
## PBYSTAND    PBYSTAND 0.82501502
## PMOTSCO     PMOTSCO 0.81362486
## MFALLEEN    MFALLEEN 0.77469168
## MBERARBO    MBERARBO 0.74527101
## PLEVEN      PLEVEN 0.73389614
## MSKB2       MSKB2 0.67226643
## MGEMOMV     MGEMOMV 0.66848679
## MSKD        MSKD 0.64817413
## MRELSA      MRELSA 0.55338948
## MOPLLAAG    MOPLLAAG 0.42198674
## MBERBOER    MBERBOER 0.40081995
## MBERZELF    MBERZELF 0.32375264
## MINK123M    MINK123M 0.29723936
## MAANTHUI    MAANTHUI 0.07381764
## ALEVEN      ALEVEN 0.02691667
## PWABEDR     PWABEDR 0.00000000
## PWALAND     PWALAND 0.00000000
## PBESAUT     PBESAUT 0.00000000
## PVRAAUT     PVRAAUT 0.00000000
## PAANHANG    PAANHANG 0.00000000

```

```
## PTRACTOR PTRACTOR 0.00000000
## PWERKT    PWERKT 0.00000000
## PBROM     PBROM 0.00000000
## PPERSONG PPERSONG 0.00000000
## PGEZONG   PGEZONG 0.00000000
## PWAOREG   PWAOREG 0.00000000
## PZEILPL   PZEILPL 0.00000000
## PPLEZIER  PPLEZIER 0.00000000
## PFIETS    PFIETS 0.00000000
## PINBOED   PINBOED 0.00000000
## AWAPART   AWAPART 0.00000000
## AWABEDR   AWABEDR 0.00000000
## AWALAND   AWALAND 0.00000000
## ABESAUT   ABESAUT 0.00000000
## AMOTSCO   AMOTSCO 0.00000000
## AVRAAUT   AVRAAUT 0.00000000
## AAANHANG  AAANHANG 0.00000000
## ATRACTOR  ATRACTOR 0.00000000
## AWERKT    AWERKT 0.00000000
## ABROM     ABROM 0.00000000
## APERSONG  APERSONG 0.00000000
## AGEZONG   AGEZONG 0.00000000
## AWAOREG   AWAOREG 0.00000000
## AZEILPL   AZEILPL 0.00000000
## APLEZIER  APLEZIER 0.00000000
## AFIETS    AFIETS 0.00000000
## AINBOED   AINBOED 0.00000000
## ABYSTAND  ABYSTAND 0.00000000
```

PPERSAUT seems to be the most important determinant of Purchase.

Part 11.c)

```
library(gmodels)
gbm_probs <- predict(gbm_fit, newdata = Caravan[test, ], n.trees = 1000, type = "response")
gbm_preds <- ifelse(gbm_probs > 0.2, 1, 0)
CrossTable(gbm_preds, Caravan$Purchase[test])
```

```
##
##
##      Cell Contents
## |-----|
## |                      N |
## | Chi-square contribution |
## |          N / Row Total |
## |          N / Col Total |
## |          N / Table Total |
## |-----|
##
##
## Total Observations in Table:  4822
##
##
```



```
##           | Caravan$Purchase[test]
##   gbm_preds |           0 |           1 | Row Total |
## -----|-----|-----|-----|
##           0 |       4359 |        252 |       4611 |
##           |       0.137 |       2.146 |           |
##           |       0.945 |       0.055 |       0.956 |
##           |       0.962 |       0.872 |           |
##           |       0.904 |       0.052 |           |
## -----|-----|-----|-----|
##           1 |        174 |         37 |        211 |
##           |       2.990 |      46.902 |           |
##           |       0.825 |       0.175 |       0.044 |
##           |       0.038 |       0.128 |           |
##           |       0.036 |       0.008 |           |
## -----|-----|-----|-----|
## Column Total |       4533 |        289 |       4822 |
##           |       0.940 |       0.060 |           |
## -----|-----|-----|-----|
##
##
```

17.5% of the people predicted to make a purchase do in fact make one. It is more accurate than the KNN and logistic predictions below.

KNN

We try KNN with $k = 2$.

```
library(class)
set.seed(1)
knn_preds <- knn(test = Caravan[test, ], train = Caravan[train, ], cl = Caravan$Purchase[train], k = 2)
CrossTable(knn_preds, Caravan$Purchase[test])
```

```
##
##
##   Cell Contents
## |-----|
## |                N |
## | Chi-square contribution |
## |      N / Row Total |
## |      N / Col Total |
## |      N / Table Total |
## |-----|
##
##
## Total Observations in Table:  4822
##
##
##           | Caravan$Purchase[test]
##   knn_preds |           0 |           1 | Row Total |
## -----|-----|-----|-----|
##           0 |       4272 |        258 |       4530 |
##           |       0.043 |       0.671 |           |
##           |       0.943 |       0.057 |       0.939 |
```

```
##           |      0.942 |      0.893 |           |
##           |      0.886 |      0.054 |           |
## -----|-----|-----|-----|
##           1 |      261 |      31 |      292 |
##           |      0.664 |     10.413 |           |
##           |      0.894 |      0.106 |      0.061 |
##           |      0.058 |      0.107 |           |
##           |      0.054 |      0.006 |           |
## -----|-----|-----|-----|
## Column Total |      4533 |      289 |      4822 |
##           |      0.940 |      0.060 |           |
## -----|-----|-----|-----|
##
##
```

10.6% of those who are predicted to purchase in fact purchase according to the KNN.

Logistic regression

We use the same threshold of 0.2 for logistic regression:

```
glm_fit <- glm(Purchase ~ ., data = Caravan[train, ], family = "binomial")
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
glm_probs <- predict(glm_fit, newdata= Caravan[test, ], type = "response")
```

```
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading
```

```
glm_preds <- ifelse(glm_probs > 0.2, 1, 0)
CrossTable(glm_preds, Caravan$Purchase[test])
```

```
##
##
##      Cell Contents
## |-----|
## |              N |
## | Chi-square contribution |
## |      N / Row Total |
## |      N / Col Total |
## |      N / Table Total |
## |-----|
##
##
## Total Observations in Table:  4822
##
##
##           | Caravan$Purchase[test]
##      glm_preds |           0 |           1 | Row Total |
## -----|-----|-----|-----|
##           0 |      4183 |      231 |      4414 |
##           |      0.271 |      4.254 |           |
##           |      0.948 |      0.052 |      0.915 |
##           |      0.923 |      0.799 |           |
##           |      0.867 |      0.048 |           |
```

```
## -----|-----|-----|-----|
##          1 |         350 |         58 |         408 |
##          |         2.934 |        46.023 |         |
##          |         0.858 |         0.142 |         0.085 |
##          |         0.077 |         0.201 |         |
##          |         0.073 |         0.012 |         |
## -----|-----|-----|-----|
## Column Total |         4533 |         289 |         4822 |
##          |         0.940 |         0.060 |         |
## -----|-----|-----|-----|
##
##
```

For logistic regression, 14.2% of those who are predicted to purchase do purchase.